

DREAMLOG

5-Phase Development Plan

From Zero to Working Prototype to Full Product

Phase 1: Foundation → Phase 2: Working Prototype → Phase 3: Launchable MVP
Phase 4: Growth Engine → Phase 5: Scale & Monetize

Version 1.0 • April 2026 • Solo Developer Build Plan

Development Philosophy

This plan is structured so that you have a usable product as early as possible and layer complexity on top. The 5 phases are designed with a clear principle: build the smallest thing that works, use it yourself, prove it matters, then expand.

Phase Map

Phase	Name	Duration	What You Have	Status
Phase 1	Foundation	Weeks 1–2	Backend + recording + transcription	Internal dev build
Phase 2	Working Prototype	Weeks 3–5	Record → AI reflection → timeline	Usable prototype (you + 10 testers)
Phase 3	Launchable MVP	Weeks 6–9	Onboarding, Mood Map, notifications, paywall	App Store ready with subscriptions
Phase 4	Growth Engine	Weeks 10–15	Pattern Radar, Life Chapters, social sharing	Retention + viral mechanics
Phase 5	Scale & Monetize	Weeks 16–24+	Dream Decoder, Relationship Map, B2B, API	Full product + revenue diversification

By the end of Phase 2, you have a working prototype you use every night. By the end of Phase 3, you have a product ready for the App Store with a subscription model. Phases 4 and 5 are about making it sticky, viral, and profitable.

PHASE 1: FOUNDATION

Timeline: Weeks 1–2 (14 days)

Goal: *Set up the entire backend, get voice recording and transcription working end-to-end. No UI polish — just the engine.*

1.1 Backend Infrastructure

Everything starts here. Build the server, database, and storage layer that will power the entire app.

1. Initialize Go project with Gin framework. Set up project structure: `cmd/api/`, `internal/handlers/`, `internal/services/`, `internal/models/`, `internal/middleware/`, `pkg/`.
2. Design and implement the PostgreSQL database schema. Core tables: `users` (`id`, `email`, `name`, `timezone`, `preferences`, `created_at`), `entries` (`id`, `user_id`, `audio_url`, `transcript`, `duration_secs`, `recorded_at`), `entry_analysis` (`entry_id`, `emotional_tone`, `topics`, `mood_score`, `reflection_text`, `summary`), `emotional_metrics` (`user_id`, `date`, `mood_score`, `dominant_emotions`).
3. Set up Cloudflare R2 (or AWS S3) bucket for audio file storage. Configure: private bucket, server-side encryption enabled, pre-signed upload URLs (15 min expiry), pre-signed download URLs for playback.
4. Set up Redis for session caching and background job queue. You'll use this for async AI processing — user records audio, uploads it, and the transcription + analysis happens in a background worker, not blocking the request.
5. Create Docker Compose file for local development: Go API server, PostgreSQL 16, Redis 7, MinIO (local S3-compatible storage for testing without cloud costs).
6. Build core API endpoints (all behind auth middleware): `POST /entries` (upload audio), `GET /entries` (list entries), `GET /entries/:id` (entry detail with analysis), `GET /me` (user profile), `PUT /me` (update preferences).
7. Write database migration scripts. Use `golang-migrate` or `goose` for version-controlled schema changes.

1.2 Authentication

1. Set up Supabase project (free tier is sufficient for MVP). Configure Google OAuth and Apple Sign-In providers.
2. Implement auth middleware in Go: validate Supabase JWT tokens on every API request, extract `user_id`, reject expired/invalid tokens.
3. Build sign-up/sign-in flow in React Native using `@supabase/supabase-js`. Support Google sign-in (both platforms), Apple sign-in (iOS), and email + magic link (fallback).
4. Implement auto-creation of user record in PostgreSQL on first sign-in (triggered by Supabase webhook or on first API call).

1.3 React Native Project Setup

1. Initialize Expo project (managed workflow). Install core dependencies: expo-av (audio recording), expo-router (navigation), @supabase/supabase-js, react-native-reanimated (animations).
2. Set up app navigation structure: Auth screens (sign-in/sign-up), Main tab navigator (Home/Record, Timeline, Insights, Settings), Entry detail screen, Recording screen (full-screen modal).
3. Configure dark theme as default (this is a nighttime app). Define color palette, typography scale, and spacing constants.
4. Set up EAS Build for iOS and Android development builds. Configure app.json with bundle identifiers, permissions (microphone access), and splash screen.

1.4 Voice Recording & Transcription Pipeline

This is the core technical pipeline. By the end of this section, a user can record their voice and get a text transcript back.

1. Build the recording screen: full-screen dark UI with one large circular record button. When recording: show live audio waveform visualization (use react-native-reanimated for smooth animation) and elapsed time counter. Add stop button and cancel button.
2. Implement audio recording using expo-av: format AAC, channels mono, sample rate 44100, bit rate 64000. This gives a ~500KB file per minute of speech — small enough for fast upload.
3. Build the upload flow: after recording stops, show a brief “Processing...” screen. Request a pre-signed upload URL from the Go API. Upload the AAC file directly to R2/S3 from the client. Send the audio URL to the Go API to trigger processing.
4. Build the transcription worker (Go background job): receive audio URL, download the file from R2, send to OpenAI Whisper API (model: whisper-1, response_format: verbose_json to get word-level timestamps and detected language). Store transcript and language in the entries table. Delete audio from R2 after successful transcription (default behavior; retention is a premium feature later).
5. Handle edge cases: accept any recording length (no minimum — even a 30-second “Today was fine” is valid). Show a gentle “Take your time” indicator at 15 minutes. Hard cap at 30 minutes. AI adapts reflection depth to match entry length (short entry = 2-sentence reflection, long entry = rich multi-paragraph analysis). Handle upload failures with retry (3 attempts with exponential backoff). Queue recordings made offline for upload when connectivity returns.
6. Test the full pipeline end-to-end: record on phone → upload → transcribe → store. Verify transcription quality across English, Hindi, and 2–3 other languages.

■ **MILESTONE:** You can record your voice on the phone and see the text transcript appear in the database. The engine works.

PHASE 2: WORKING PROTOTYPE

Timeline: Weeks 3–5 (21 days)

Goal: Add the AI brain: reflections that reference your history, a browsable timeline of past entries, and basic mood tracking. By the end of this phase, you use the app every night and it feels valuable.

2.1 AI Analysis Engine

This is where DreamLog becomes intelligent. The AI reads your transcript, understands the emotional content, and generates a personalized reflection.

1. Build the Claude API integration service in Go. Use Claude Sonnet for all nightly analysis (best balance of quality and speed). Structure: send transcript + context → receive structured JSON response with fields: `emotional_tone` (array of emotions with intensity 0–1), `topics` (array of key subjects discussed), `mood_score` (1–100), `key_quotes` (2–3 notable phrases from the user), `summary` (2–3 sentence summary of the entry), `reflection` (3–5 sentence personalized reflection).
2. Design and iterate on the master analysis prompt. This is the most important piece of engineering in the entire product. The prompt must instruct Claude to: analyze the emotional content, identify key themes, reference the user's previous entries and patterns (provided as context), write a reflection that feels warm but not patronizing, insightful but not presumptuous. Never diagnose, never prescribe. Ask exactly one thought-provoking question at the end. Spend real time on this — the quality of this prompt determines whether users come back.
3. Implement entry summarization: for every entry, generate a 2–3 sentence summary. Store this alongside the full transcript. These summaries become the building blocks for giving the AI context about the user's history without exceeding token limits.
4. Build the context assembly pipeline. When a new entry comes in, assemble the context package sent to Claude: the user's profile (name, how long they've been using the app), summaries of the last 5 entries, the full transcript of the last entry (for continuity), and any notable patterns or themes that have emerged. Keep total context under 4,000 tokens at this stage (optimize later). This is what makes reflections feel personal — the AI knows what you said yesterday and last week.
5. Implement crisis detection. Before generating the reflection, run a separate screening pass on the transcript. Check for: mentions of self-harm, suicidal ideation, harm to others, severe substance abuse crisis. If detected: skip the normal reflection. Instead, show a compassionate message acknowledging their pain and display localized crisis resources (988 Lifeline in US, iCall/Vandrevala in India, etc.). This is non-negotiable and must be present from the first version anyone uses.
6. Build the processing pipeline end-to-end: audio uploaded → Whisper transcription → crisis screening → context assembly → Claude analysis → store analysis → push notification to user that reflection is ready. Total processing time target: under 15 seconds.

2.2 Reflection Display & Goodnight Flow

1. Build the reflection screen. After the recording is processed (typically 10–15 seconds), the user sees their personalized reflection. Design: dark calming background, reflection text in a clean readable font, gentle fade-in animation. Show the mood score as a subtle color indicator (not a number — keep it emotional, not clinical).
2. Add two action buttons below the reflection: “Goodnight” (closes the app with a gentle animation and optional sleep sound) and “Tell me more” (opens a brief 2–3 exchange follow-up conversation).
3. Implement the “Tell me more” flow. This is a short, focused conversation — not an open-ended chatbot. The AI continues from the reflection’s closing question. User replies via voice (preferred) or text. AI responds with one more insight or question. After 2–3 exchanges, the AI gently wraps up: “Something to sit with tonight. Sleep well.”
4. Build the “processing” waiting screen shown while the AI analyzes. Show a calming animation (breathing circle, gentle pulse, stars) — not a loading spinner. The wait should feel intentional, like the AI is “thinking about what you said.”

2.3 Entry Timeline

1. Build the timeline screen: a scrollable list of all past entries, newest first. Each entry card shows: date and time, mood indicator (color dot or emoji derived from mood_score), 1-line preview of the summary, topics as small tags.
2. Build the entry detail screen. Tap any entry card to see: full transcript (scrollable), the AI reflection from that night, mood score, emotional tones detected, and topics. If the user opted for the “Tell me more” conversation, show that exchange too.
3. Implement basic keyword search across all entry transcripts. Simple full-text search using PostgreSQL’s built-in tsvector/tsquery. Search bar at the top of the timeline.

2.4 Basic Mood Tracking

1. Store daily mood_score (1–100) from each entry’s AI analysis in the emotional_metrics table.
2. Build a simple 7-day mood chart on the home screen. A small line graph showing the last 7 days of mood scores. Tap any day to jump to that entry. This gives the user a visual sense of their emotional week at a glance.
3. Add a daily “streak” counter on the home screen: “5-day streak 🌟”. Simple but surprisingly effective for building the daily habit.

2.5 Morning Nudge (Basic)

1. After each nightly analysis, generate a morning nudge: a single personalized sentence to send as a push notification the next morning. Store in the database with a scheduled delivery time (user’s timezone + their preferred morning time, default 7:30am).
2. Set up Firebase Cloud Messaging (FCM) in the React Native app. Request notification permissions during onboarding.

3. Build a simple notification scheduler in Go: a cron job that runs every minute, checks for nudges due for delivery, sends them via FCM. Mark as sent after delivery.

▮ **MILESTONE:** You have a working prototype. Every night: open app → talk for 2 minutes → see a personalized reflection that references your past entries → wake up to a morning nudge. You browse your timeline and see your mood trend. Share with 10 friends/testers for feedback.

PHASE 3: LAUNCHABLE MVP

Timeline: Weeks 6–9 (28 days)

Goal: *Polish the UX, add onboarding, build the full Mood Map, implement subscriptions via RevenueCat, and prepare for App Store submission. By the end of this phase, real users can download, use, and pay for DreamLog.*

3.1 Onboarding Flow

First impression determines whether users stay. The onboarding must be fast (under 90 seconds), warm, and lead to an immediate “aha” moment with their first reflection.

1. Screen 1 — Welcome: “DreamLog is your nightly companion. Talk, and it listens. Over time, it’ll understand you better than you understand yourself.” Calming animation. “Get Started” button.
2. Screen 2 — Sign In: Google / Apple / email. One tap, no friction.
3. Screen 3 — Quick Setup: Name (pre-filled from auth if available), timezone (auto-detected), usual bedtime (time picker — used for reminder scheduling), preferred morning nudge time (default 7:30am).
4. Screen 4 — Emotional Baseline: “In one sentence, how would you describe where you are in life right now?” Voice or text input. This seeds the AI’s initial context so even the first reflection feels personal.
5. Screen 5 — First Recording: “Take a moment. How was your day? What’s on your mind tonight?” Soft prompt to reduce blank-canvas anxiety. Record button prominent.
6. Screen 6 — First Reflection: Show the AI’s reflection for their first entry. THIS is the make-or-break moment. If the user feels understood, they’ll come back tomorrow. The AI must reference their baseline statement and entry to feel personal even on Day 1.
7. Screen 7 — Habit Setup: “We’ll remind you each evening and send you a thought each morning.” Confirm notification permissions. Explain 7-day free trial of Pro. Done.

3.2 Full Mood Map

1. Extend mood tracking beyond 7 days. Store and display the full emotional timeline from the user’s first entry.
2. Build the Mood Map screen: an interactive line chart that shows mood_score over time. Pinch to zoom between day/week/month/all-time views. Tap any data point to see that day’s entry summary in a bottom sheet.
3. Add emotion tags below the chart: show the dominant emotions for the visible time period as colored tags (e.g., “anxious” in yellow, “hopeful” in green, “frustrated” in red). This makes the mood data feel human, not just a number.
4. Build correlation hints: simple overlays showing day-of-week patterns. “Your mood tends to dip on Mondays and peak on Fridays.” These are computed server-side after 3+ weeks of data.

5. Add exportable mood summary (plain text or screenshot) that users can share with a therapist. “Share my last 30 days” button generates a clean image.

3.3 Notification System (Full)

1. Build the complete notification system with all notification types: evening reminder (30 min before bedtime), morning nudge (personalized to last entry), streak-at-risk (if no entry for 2 days — gentle, never guilt-inducing), weekly report ready (Sunday evening), milestone celebrations (7, 30, 90, 365 days).
2. Implement smart evening reminders that vary nightly. Rotate between: a question (“Anything worth talking about today?”), a gentle nudge (“Your DreamLog is ready when you are”), a pattern tease (“I noticed something about your week...”). Repetitive notifications get ignored; varied ones get opened.
3. Build notification preferences screen: toggle each type on/off, set custom times for evening and morning notifications, choose quiet days (e.g., no reminders on weekends).

3.4 Context Window Optimization

As users accumulate entries, you need a smarter system for giving the AI relevant context without exceeding token limits or API costs.

1. Implement weekly digest generation: every Sunday, a background job compresses that week’s entry summaries into a single “weekly digest” (3–5 sentences capturing the week’s emotional arc, key themes, and notable moments). Store in a `weekly_digests` table.
2. Implement monthly profile generation: at the end of each month, compress weekly digests into a “monthly emotional profile” (key patterns, growth areas, recurring themes, relationship dynamics). Store in `monthly_profiles` table.
3. Update the context assembly pipeline: for nightly reflections, send Claude: user profile + current month’s weekly digests + last 3 entry summaries + last full entry transcript + any active patterns. Target: under 6K tokens total context, giving Claude room for a quality response.
4. Set up pgvector extension on PostgreSQL. Generate embeddings for each entry summary (using OpenAI embeddings API or a smaller model). When assembling context, find the 2 most semantically similar past entries to tonight’s transcript and include their summaries. This allows the AI to reference relevant past entries even from months ago.

3.5 Subscription System

1. Set up RevenueCat account. Configure products in App Store Connect and Google Play Console: `dreamlog_pro_monthly` (\$7.99), `dreamlog_pro_annual` (\$59.99), `dreamlog_premium_monthly` (\$14.99), `dreamlog_premium_annual` (\$99.99).

2. Integrate RevenueCat SDK in React Native. Implement: check subscription status on app launch, cache status locally for offline access, listen for subscription changes.
3. Build the paywall screen. Show a clear comparison of Free vs Pro vs Premium features (use the tier table from the PRD). Include: price, what's included, annual savings calculation, "Start 7-day free trial" button for Pro. Make it clean, not aggressive.
4. Implement smart trial trigger: the 7-day Pro trial prompt appears after the user's 2nd entry (not at sign-up). This ensures the user has experienced the core loop and seen two reflections before being asked to consider paying. Much higher conversion than prompting cold.
5. Build feature gating throughout the app. Free tier: 5 entries/month, basic reflection (no history references), 7-day mood chart only, no Pattern Radar, no deep dive conversation. Pro tier: unlimited entries, deep reflections with full history, full Mood Map, Pattern Radar (Phase 4), weekly reports, 3 deep dive conversations/week. Premium tier: everything in Pro plus Dream Decoder, Relationship Map, unlimited deep dives, therapist reports, audio retention.
6. Implement the trial-end conversion screen. When the 7-day trial ends, show a personalized summary: "In 7 days, DreamLog tracked your mood across X entries, learned about [topics], and gave you [N] personalized reflections. Want to keep going?" Include the mood chart from their trial period. This converts significantly better than a generic "Trial ended" message.
7. Set up regional pricing in RevenueCat: India/SEA/LATAM at approximately 40% of US price. Student discount at 50% off with .edu verification (can be manual initially).
8. Build subscription management screen: view current plan, switch plans, cancel subscription, restore purchases. Make cancellation easy and non-guilt-inducing — trust builds long-term retention better than dark patterns.

3.6 App Store Preparation

1. Design App Store screenshots (6 screens): recording screen, reflection example, mood map, pattern tease, morning nudge example, "your AI listener" headline.
2. Record App Store preview video (30 seconds): show the core loop — open app, record, get reflection, see mood map.
3. Write App Store listing: title, subtitle ("Your nightly AI listener"), description, keywords. Optimize for: journaling, mental health, mood tracker, AI journal, voice journal.
4. Add required legal: privacy policy (detail data handling, audio deletion policy, encryption), terms of service, mental health disclaimer ("DreamLog is a wellness companion, not a replacement for professional care").
5. Implement "Get Help Now" button accessible from all screens — links to crisis hotlines. Required for App Store approval in the health/wellness category.
6. Submit to both App Store and Google Play. Plan for 1–2 weeks of review time. Address any rejection feedback immediately.

▮ **MILESTONE:** DreamLog is LIVE on the App Store and Google Play. Users can download, sign up, use the free tier, start a Pro trial, and subscribe. You have a real product generating real revenue.

PHASE 4: GROWTH ENGINE

Timeline: Weeks 10–15 (42 days)

Goal: Add the features that make DreamLog genuinely irreplaceable — Pattern Radar, Life Chapters, weekly reports, and viral sharing mechanics. This phase drives retention and organic growth.

4.1 Pattern Radar

This is the single most differentiating feature in the product. After 2–3 weeks of entries, the AI starts surfacing emotional patterns the user cannot see themselves.

1. Build the pattern detection pipeline. Runs as a weekly background job for each user with 10+ entries. Sends the user's full history (monthly profiles + weekly digests + recent entry summaries) to Claude Opus (use the more powerful model for this task) with a specialized prompt asking it to identify patterns across 6 categories:
 - Recurring triggers: "You feel most anxious on Sunday evenings — 6 of the last 8 Sundays"
 - People patterns: "You consistently feel drained after interactions with [person]"
 - Contradictions: "You say you're over the breakup but you've mentioned your ex in 9 of 12 entries"
 - Avoidance: "You've mentioned wanting to confront your manager 7 times but always change the subject"
 - Growth moments: "One month ago you were paralyzed about the move. This week you're actively planning it."
 - Cyclical: "Your mood dips in the last week of every month — could be deadline-related or hormonal"
2. Store detected patterns in the patterns table with: type, description, confidence score, evidence (array of entry IDs that support it), detected_at, status (new/seen/dismissed).
3. Build the Pattern Radar screen: card-based feed of detected patterns. Each card shows: pattern type icon, description, "Evidence: based on X entries over Y weeks", and a "See entries" link to the supporting evidence. New patterns have a glowing indicator.
4. Add pattern push notification: "I found a new pattern in your recent entries." Pro/Premium users only. This notification has the highest open rate because it triggers curiosity.
5. Implement pattern feedback: user can mark each pattern as "Spot on" or "Not quite right." Feed this back into future pattern detection prompts to improve accuracy over time.
6. Build the locked pattern teaser for free users. After they accumulate enough entries, show: "We've detected 2 patterns in your entries" with a blurred card. Tap to see the upgrade prompt. This is the highest-converting paywall trigger.

4.2 Life Chapters

1. Build the chapter detection algorithm. Use entry embeddings (stored in pgvector) to cluster entries by emotional themes and topics. When a cluster of 5+ entries shares a dominant theme, create a Life Chapter.
2. Auto-generate chapter titles and summaries using Claude. Examples: “The Job Transition” (March–May: 47 entries, dominant emotion: anxiety → excitement), “The Friendship Shift” (April: 12 entries about evolving feelings).
3. Build the Life Chapters screen: visual timeline showing chapters as segments with mood arcs. Tap a chapter to see: summary, mood trajectory within the chapter, key turning points, and entry list.
4. Implement auto-closing of chapters: when the AI detects the user has moved on from a theme (no related entries for 2+ weeks), the chapter is marked as closed with a wrap-up summary.
5. Premium feature: allow users to manually create chapters and tag entries. Some users will want to organize their emotional life themselves.

4.3 Weekly Reports

1. Build the weekly report generator. Runs every Sunday evening. Analyzes all entries from the past 7 days and produces: emotional summary of the week (3–4 sentences), mood trend visualization (mini chart), top 3 themes, notable moments (specific quotes or insights), one pattern observation, and a “focus for next week” suggestion.
2. Design the weekly report screen: a beautiful, scrollable, card-based report that feels like a personal newsletter about your inner life. Use warm colors, clean typography, and the user's own words quoted back to them.
3. Build a shareable report card: a single image (optimized for Instagram Stories format, 1080x1920) showing the week's mood arc, dominant emotion, and one sanitized insight. Users can share without revealing private details. This is the primary organic growth driver.
4. Send weekly report notification on Sunday evening. Email delivery option for users who want it in their inbox (for re-engagement of lapsed users).

4.4 Sharing & Viral Mechanics

1. Build the sharing infrastructure: allow users to share their weekly report card, milestone achievements (30-day streak, first chapter completed), and Mood Map screenshots with branded DreamLog overlay.
2. Implement referral program: user shares a unique invite link. Friend gets 2 weeks of Pro free. User gets 2 weeks of Pro free when friend signs up. Track via RevenueCat promotional offers.
3. Create a “30 Days of DreamLog” challenge. In-app prompt for new users: “Commit to 30 consecutive nights. See what you learn about yourself.” Milestone cards at day 7, 14, 21, 30 — each shareable.

4. Build in-app “Invite a friend” screen accessible from settings and from milestone celebrations.

4.5 Marketing Push

- Launch social media content strategy: TikTok and Instagram Reels. Content pillars: “My AI noticed a pattern I couldn’t see” (pattern reveals), “What DreamLog told me after 30 days” (transformation stories), “The pattern that changed my life” (emotional hooks). Post 3–5 times/week.
- Product Hunt launch: prepare compelling copy, demo video, and maker story. Target a Monday launch for maximum visibility.
- Reddit seeding: authentic posts in r/selfimprovement, r/mentalhealth, r/journaling, r/productivity, r/getdisciplined. Share your own experience using it.
- Reach out to 20–50 wellness and self-improvement influencers for beta access and honest reviews.
- SEO content: publish articles on “voice journaling benefits”, “how to track your mental health”, “why you can’t stick to journaling”, “AI therapy alternatives.”
- Build a simple landing page / marketing website with: hero video, feature showcase, testimonials from beta users, download links, and email capture for blog updates.

■ **MILESTONE:** DreamLog has Pattern Radar, Life Chapters, weekly reports, and viral sharing. Users who stay past 2 weeks rarely churn because the data moat is too valuable. Organic growth is ramping from social sharing. Target: \$5K–\$10K MRR.

PHASE 5: SCALE & MONETIZE

Timeline: Weeks 16–24+ (Ongoing)

Goal: Add premium differentiators (*Dream Decoder*, *Relationship Map*, *Therapist Prep*), build additional revenue streams (*Guided Journeys*, *Year in Review*, *B2B*), and optimize for profitability.

5.1 Dream Decoder

1. Build dream detection: when the AI analyzes a nightly entry, check if the user described a dream (mention of “I dreamed”, “I had a dream”, or vivid narrative that sounds like dream content).
2. Implement dream analysis prompt: uses established psychological frameworks (Jungian archetypes, common dream symbolism research) to interpret the dream. Critically, it connects dream themes to the user’s recent emotional patterns. “You dreamed about being lost in a building. You’ve been talking about feeling directionless at work for two weeks.”
3. Build dream journal section: separate tab or section in the timeline filtering only dream entries. Each shows: the dream narrative, the interpretation, and the connection to waking emotional themes.
4. Premium-only feature. Free/Pro users see: “Dream detected! Upgrade to Premium for dream analysis.”

5.2 Relationship Map

1. Build person extraction from transcripts. Use Claude to detect: names, nicknames (“my boss”, “mom”, “my ex”), and the emotional context in which they’re mentioned. Store in `person_mentions` table with sentiment score per mention.
2. Build the Relationship Map visualization: an interactive graph where each person is a node. Node size = mention frequency. Node color = average sentiment (green = positive, yellow = mixed, red = negative). Tap a person to see their detail view.
3. Build per-person detail view: sentiment history over time (chart), notable quotes about them, and an AI-generated relationship insight (“You give a lot of emotional energy to [person] but rarely describe receiving support back”).
4. Implement “missing connection” alerts: if a frequently mentioned person disappears from entries for 3+ weeks, note it. “You haven’t mentioned [friend] in a month. You used to talk about them often.”
5. Premium-only feature.

5.3 Therapist Prep Report

1. Build the therapist prep report generator: user taps “Prepare for therapy session” and selects the date range (last week, last 2 weeks, since last session). The AI generates: top 3 themes to discuss, emotional trend summary, specific moments worth exploring, detected patterns, and suggested questions for the therapist.

2. Export as a clean PDF with DreamLog branding. User can email it to their therapist or read from their phone during the session.
3. Build a therapist-facing landing page: explain DreamLog to therapists, show a sample prep report, encourage them to recommend it to clients. This becomes a high-trust acquisition channel.

5.4 Guided Journeys (Premium Content)

1. Build the Journey framework: a multi-week structured experience with specific nightly prompts, tailored analysis, and milestone reflections. Users opt into a journey and receive a specific prompt each night instead of the open-ended default.
2. Launch initial journeys: “Processing Grief” (8 weeks), “Career Clarity” (6 weeks), “Relationship Audit” (4 weeks), “Anxiety Toolkit” (4 weeks), “New Parent Adjustment” (6 weeks). Each designed with input from psychological frameworks.
3. Monetize as one-time purchases (\$4.99 each) or included in Premium tier. Users can run one journey at a time alongside their regular nightly check-ins.

5.5 Year in Review

1. Build the annual review generator: a comprehensive AI-written narrative of the user’s year. Includes: life chapters timeline, emotional arc visualization, top patterns discovered, growth milestones, relationship shifts, and a personalized letter from the AI reflecting on their journey.
2. Offer as a beautifully designed digital PDF (\$9.99) or a printed hardcover book shipped to the user (\$29.99 via a print-on-demand partner like Lulu or Blurb).
3. This becomes DreamLog’s “Spotify Wrapped” moment — an annual event that drives massive social sharing and re-engagement. Launch in December for maximum impact.

5.6 Voice Tone Analysis

1. Build a Python microservice using librosa or pyAudioAnalysis that analyzes the raw audio (before deletion) for: speaking pace (words per minute), pitch variability, energy levels, pause frequency, and vocal tension. These voice-level signals add a layer of emotional data that transcription alone misses.
2. Feed voice metrics into the mood scoring model: someone who speaks quickly with short pauses and high pitch is likely anxious, even if their words sound calm.
3. Display voice insights in the entry detail: “Your voice was notably quieter and slower tonight compared to your usual pace.”

5.7 Enterprise / B2B Wellness

1. Build an anonymized team mood dashboard for HR/people teams. Companies purchase DreamLog as an employee wellness benefit at \$5/employee/month. Employees get free Pro access; HR sees only aggregate, anonymized mood trends — never individual data.
2. Dashboard shows: team-level mood trends over time, stress spikes correlated with company events (quarter-end, reorgs, layoffs), engagement scores, and anonymous sentiment themes.
3. This revenue stream de-risks the business from pure consumer subscription dependency and commands higher per-user pricing.

5.8 Platform & API

- Build a public REST API for the emotional analysis engine: send text, get back emotional tone, mood score, topics, and patterns. License to other wellness apps, therapy platforms, and HR tools.
- Apple Watch companion app: quick 30-second voice check-in from your wrist, haptic morning nudge.
- Integration with Apple Health (mood tracking category) and Google Health Connect.
- Explore Alexa/Google Home integration: “Hey Google, talk to DreamLog” for a fully hands-free bedtime check-in.

■ **MILESTONE:** DreamLog is a full-featured mental wellness platform with multiple revenue streams: consumer subscriptions, premium content (Journeys), annual products (Year in Review), B2B licensing, and API access. Target: \$30K–50K+ MRR.

Complete Task Reference

Every task across all 5 phases. Priority: P0 = required for that phase, P1 = high value, P2 = important but can ship without. Status: Proto = needed for prototype (Phase 2), MVP = needed for App Store launch (Phase 3), Growth/Scale = Phase 4/5.

#	Task	Pri	Effort	Gate	Notes
1.1	Go backend + Gin + API scaffolding	P0	3 hrs	Proto	Foundation for everything
1.2	PostgreSQL schema (users, entries, analysis, metrics)	P0	5 hrs	Proto	Design for future extensibility
1.3	R2/S3 encrypted audio storage + signed URLs	P0	3 hrs	Proto	Privacy from Day 1
1.4	Redis cache + background job queue	P0	2 hrs	Proto	Async AI processing
1.5	Docker Compose local dev environment	P1	2 hrs	Proto	Go + PG + Redis + MinIO
1.6	Core API endpoints (CRUD entries + user)	P0	4 hrs	Proto	REST API foundation
1.7	DB migration scripts (golang-migrate)	P1	2 hrs	Proto	Version-controlled schema
1.8	Supabase Auth setup + JWT middleware	P0	4 hrs	Proto	Google + Apple + email sign-in
1.9	React Native Expo project init + nav	P0	3 hrs	Proto	Tabs, screens, dark theme
1.10	EAS Build config (iOS + Android)	P1	2 hrs	Proto	Dev builds for testing
1.11	Recording screen + waveform visualization	P0	6 hrs	Proto	Core UX — make it beautiful
1.12	Audio recording (expo-av, AAC mono 64kbps)	P0	4 hrs	Proto	Test across 5+ devices
1.13	Background upload with retry + offline queue	P0	4 hrs	Proto	Reliability is non-negotiable
1.14	Whisper API transcription integration	P0	4 hrs	Proto	Language auto-detect
1.15	Full transcription pipeline (upload → transcribe → store → delete)	P0	4 hrs	Proto	Privacy-first by default
1.16	Edge cases (silence, noise, short audio)	P1	3 hrs	Proto	Graceful error handling
2.1	Claude API integration + structured JSON output	P0	6 hrs	Proto	Core intelligence layer
2.2	Master analysis prompt (tone, topics, mood, reflection)	P0	10 hrs	Proto	Most important engineering work
2.3	Entry summarization (2–3 sentences per entry)	P0	3 hrs	Proto	Building block for context
2.4	Context assembly pipeline (history-aware)	P0	6 hrs	Proto	What makes reflections personal
2.5	Crisis detection screening	P0	6 hrs	Proto	Non-negotiable safety feature
2.6	Full processing pipeline (audio → transcript → screen → analyze → notify)	P0	4 hrs	Proto	End-to-end pipeline
2.7	Reflection display screen + Goodnight/Tell me more	P0	5 hrs	Proto	The magic moment UI
2.8	Tell me more flow (2–3 voice/text exchanges)	P1	5 hrs	Proto	Optional depth after reflection
2.9	Processing waiting screen (calming animation)	P1	2 hrs	Proto	Make the wait feel intentional

2.10	Entry timeline (scrollable list + cards)	P0	4 hrs	Proto	Browse past entries
2.11	Entry detail screen (transcript + reflection)	P0	3 hrs	Proto	Full entry view
2.12	Keyword search across entries	P2	3 hrs	Proto	PostgreSQL full-text search
2.13	7-day mood chart on home screen	P1	3 hrs	Proto	Visual emotional snapshot
2.14	Streak counter on home screen	P1	1 hr	Proto	Habit motivation
2.15	Morning nudge generation + scheduling	P0	5 hrs	Proto	Personalized next-morning push
2.16	FCM push notification integration	P0	3 hrs	Proto	Delivery infrastructure
2.17	Basic notification scheduler (cron)	P0	3 hrs	Proto	Morning nudge + evening reminder
3.1	Onboarding flow (7 screens)	P0	8 hrs	MVP	First impression is everything
3.2	Full Mood Map (interactive, zoomable, all-time)	P0	8 hrs	MVP	Major feature for marketing
3.3	Emotion tags on Mood Map	P1	3 hrs	MVP	Human-readable mood data
3.4	Day-of-week correlation hints	P2	3 hrs	MVP	Simple pattern preview
3.5	Exportable mood summary for therapists	P2	2 hrs	MVP	Screenshot + share
3.6	Full notification system (5 types + variance)	P0	6 hrs	MVP	Retention driver
3.7	Notification preferences screen	P1	3 hrs	MVP	User control
3.8	Weekly digest generation (background job)	P0	4 hrs	MVP	Context window management
3.9	Monthly profile generation	P1	3 hrs	MVP	Long-term memory compression
3.10	pgvector setup + entry embeddings	P1	4 hrs	MVP	Semantic retrieval for old entries
3.11	Updated context assembly with digests + vectors	P0	4 hrs	MVP	Smarter, cheaper AI calls
3.12	RevenueCat integration (iOS + Android)	P0	6 hrs	MVP	Payment infrastructure
3.13	Paywall screen (Free/Pro/Premium comparison)	P0	4 hrs	MVP	Conversion UI
3.14	7-day trial (triggered after 2nd entry)	P0	3 hrs	MVP	Smart trial timing
3.15	Feature gating system	P0	5 hrs	MVP	Lock features by tier
3.16	Trial-end conversion screen (personalized)	P0	4 hrs	MVP	Highest-converting moment
3.17	Regional pricing + student discount	P1	3 hrs	MVP	Expand market
3.18	Subscription management screen	P0	3 hrs	MVP	Upgrade/cancel/restore
3.19	App Store screenshots + preview video	P0	4 hrs	MVP	ASO critical
3.20	App Store listing + keywords	P0	2 hrs	MVP	Discoverability
3.21	Privacy policy + terms + disclaimers	P0	3 hrs	MVP	Legal requirement
3.22	Get Help Now button + crisis resources	P0	2 hrs	MVP	App Store requirement
3.23	Submit to App Store + Google Play	P0	2 hrs	MVP	Go live
4.1	Pattern detection pipeline (6 types)	P0	12 hrs	Growth	Crown jewel feature
4.2	Pattern storage + evidence linking	P0	3 hrs	Growth	Data layer for patterns
4.3	Pattern Radar UI (cards + evidence)	P0	6 hrs	Growth	Compelling display
4.4	Locked pattern teaser for free users	P0	2 hrs	Growth	Highest-converting paywall

4.5	Pattern notifications + feedback	P1	3 hrs	Growth	Engagement + accuracy
4.6	Chapter detection algorithm (embedding clusters)	P1	8 hrs	Growth	AI narrative segmentation
4.7	Chapter auto-titling + summaries	P1	4 hrs	Growth	AI-generated story
4.8	Life Chapters screen + chapter detail	P1	6 hrs	Growth	Visual life story
4.9	Weekly report generator	P0	8 hrs	Growth	Sunday retention event
4.10	Shareable report card (Stories format)	P0	4 hrs	Growth	Organic growth driver
4.11	Weekly report email delivery	P1	3 hrs	Growth	Re-engagement channel
4.12	Sharing infrastructure + branded overlays	P1	4 hrs	Growth	Social sharing
4.13	Referral program (2 weeks free each)	P1	4 hrs	Growth	Viral acquisition
4.14	30-Day Challenge feature	P2	4 hrs	Growth	Engagement campaign
4.15	Landing page / marketing website	P0	6 hrs	Growth	SEO + conversion
4.16	Social media content strategy	P0	Ongoing	Growth	TikTok, Reels, Reddit
4.17	Product Hunt launch prep	P1	4 hrs	Growth	One-time event
5.1	Dream detection from transcripts	P1	4 hrs	Scale	Premium differentiator
5.2	Dream analysis prompt + UI	P1	6 hrs	Scale	Fascinating and shareable
5.3	Dream journal section	P2	3 hrs	Scale	Filtered dream timeline
5.4	Person mention extraction (NLP)	P1	6 hrs	Scale	Names + sentiment
5.5	Relationship Map visualization	P1	8 hrs	Scale	Interactive graph
5.6	Per-person detail + insights	P2	4 hrs	Scale	Relationship intelligence
5.7	Missing connection alerts	P2	2 hrs	Scale	Proactive relationship care
5.8	Therapist prep report generator	P1	6 hrs	Scale	PDF export
5.9	Therapist landing page	P2	4 hrs	Scale	Professional referral channel
5.10	Guided Journey framework	P1	8 hrs	Scale	Multi-week structured experiences
5.11	3–5 initial Journeys (content + prompts)	P1	10 hrs	Scale	Premium content revenue
5.12	Year in Review generator	P2	10 hrs	Scale	Annual Spotify Wrapped moment
5.13	Voice tone analysis microservice	P2	8 hrs	Scale	Audio-level emotion signals
5.14	B2B team mood dashboard	P2	16 hrs	Scale	Enterprise revenue stream
5.15	Public API for emotional analysis	P3	12 hrs	Scale	Platform licensing
5.16	Apple Watch companion	P3	8 hrs	Scale	Wrist-level check-ins

Timeline Summary

Week	Phase	What Gets Done	Users	Revenue
1–2	Phase 1	Backend, auth, recording, transcription pipeline	Just you	\$0
3–4	Phase 2a	AI reflections, reflection screen, timeline	You (daily)	\$0
5	Phase 2b	Mood chart, morning nudge, search	10 testers	\$0
6–7	Phase 3a	Onboarding, full Mood Map, notifications	50 beta	\$0
8–9	Phase 3b	Subscriptions, App Store prep, launch	500+	\$500 MRR
10–12	Phase 4a	Pattern Radar, Life Chapters, weekly reports	2,000+	\$3K MRR
13–15	Phase 4b	Sharing, referrals, marketing push	5,000+	\$10K MRR
16–20	Phase 5a	Dream Decoder, Relationship Map, Journeys	10,000+	\$20K MRR
21–24+	Phase 5b	Year in Review, B2B, API, Watch app	20,000+	\$50K+ MRR

Phase 1–2: Make it work.

Phase 3: Make it launchable.

Phase 4: Make it grow.

Phase 5: Make it a business.

Start tonight. Record your first entry. Build what you need.